

Introduction to Machine Learning Lab

February 16, 2026

1 Lab 01: Introduction to Machine Learning and Python Environment

Importing Libraries

```
[1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import warnings  
warnings.filterwarnings('ignore')
```

```
[2]: from sklearn.datasets import load_iris
```

Loading Iris Dataset

```
[3]: iris = load_iris()
```

```
[4]: df = pd.DataFrame(iris.data, columns=iris.feature_names)  
df["species"] = iris.target
```

Displaying basic statistics and first few records

```
[5]: df.head()
```

```
[5]:    sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  \\\n0                 5.1           3.5            1.4            0.2  
1                 4.9           3.0            1.4            0.2  
2                 4.7           3.2            1.3            0.2  
3                 4.6           3.1            1.5            0.2  
4                 5.0           3.6            1.4            0.2  
  
  species  
0      0  
1      0  
2      0  
3      0  
4      0
```


```
[6]: df.describe()
```

```
[6]:      sepal length (cm)  sepal width (cm)  petal length (cm) \
count          150.000000        150.000000        150.000000
mean           5.843333         3.057333         3.758000
std            0.828066         0.435866         1.765298
min            4.300000         2.000000         1.000000
25%           5.100000         2.800000         1.600000
50%           5.800000         3.000000         4.350000
75%           6.400000         3.300000         5.100000
max            7.900000         4.400000         6.900000


      petal width (cm)  species
count          150.000000        150.000000
mean           1.199333         1.000000
std            0.762238         0.819232
min            0.100000         0.000000
25%           0.300000         0.000000
50%           1.300000         1.000000
75%           1.800000         2.000000
max            2.500000         2.000000
```

Plotting histogram and scatter using matplotlib

```
[7]: plt.hist(df["sepal length (cm)"])
plt.title("Distribution of Sepal Length")
plt.xlabel("Sepal Length")
plt.ylabel("Frequency")
plt.show()
```




```
[8]: plt.scatter(df["sepal length (cm)"], df["petal length (cm)"])
plt.xlabel("Sepal Length")
plt.ylabel("Petal Length")
plt.title("Sepal vs Petal Length")
plt.show()
```



Finding correlation and showing heatmap using seaborn

```
[9]: corr = df.corr()
sns.heatmap(corr, annot=True)
plt.title("Feature Correlation")
plt.show()
```



Splitting data into test and train set

```
[10]: from sklearn.model_selection import train_test_split

X = df.drop("species", axis=1)
y = df["species"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
    ↪2, random_state=42)
```

```
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (120, 4)
X_test shape: (30, 4)
y_train shape: (120,)
```

```
y_test shape: (30,)
```

2 Lab – 02: Data Preprocessing

Loading Titanic Dataset

```
[11]: df = sns.load_dataset("titanic")
df.head()
```

```
[11]:    survived  pclass      sex   age  sibsp  parch     fare embarked  class \
0          0       3    male  22.0      1      0    7.2500      S  Third
1          1       1  female  38.0      1      0   71.2833      C  First
2          1       3  female  26.0      0      0    7.9250      S  Third
3          1       1  female  35.0      1      0   53.1000      S  First
4          0       3    male  35.0      0      0    8.0500      S  Third

      who  adult_male  deck  embark_town  alive  alone
0  man        True    NaN  Southampton    no  False
1 woman      False     C  Cherbourg    yes  False
2 woman      False    NaN  Southampton    yes   True
3 woman      False     C  Southampton    yes  False
4  man        True    NaN  Southampton    no   True
```

Checking features null values

```
[12]: df.isnull().sum()
```

```
[12]: survived      0
pclass         0
sex           0
age        177
sibsp         0
parch         0
fare          0
embarked      2
class         0
who           0
adult_male     0
deck        688
embark_town    2
alive         0
alone         0
dtype: int64
```

Dropping Null values

```
[13]: df_drop = df.dropna()
print("Original shape:", df.shape)
```

```
print("After drop:", df_drop.shape)
```

Original shape: (891, 15)

After drop: (182, 15)

Filling Null values with median

```
[14]: df["age"].fillna(df["age"].median(), inplace=True)
```

Identifying categorical columns.

```
[15]: df.select_dtypes(include=["object", "category"]).head()
```

```
[15]:   sex embarked class who deck embark_town alive
  0 male         S Third man   NaN Southampton no
  1 female       C First woman   C Cherbourg yes
  2 female       S Third woman  NaN Southampton yes
  3 female       S First woman   C Southampton yes
  4 male         S Third man   NaN Southampton no
```

Applying Label Encoder

```
[16]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df["class_encoded"] = le.fit_transform(df["class"])
df.head()
```

```
[16]:    survived pclass      sex   age  sibsp  parch      fare embarked class \
  0        0      3    male  22.0     1      0    7.2500      S Third
  1        1      1  female  38.0     1      0   71.2833      C First
  2        1      3  female  26.0     0      0    7.9250      S Third
  3        1      1  female  35.0     1      0   53.1000      S First
  4        0      3    male  35.0     0      0    8.0500      S Third

      who  adult_male deck  embark_town alive  alone  class_encoded
  0  man      True  NaN  Southampton  no  False          2
  1 woman     False   C  Cherbourg  yes  False          0
  2 woman     False  NaN  Southampton  yes   True          2
  3 woman     False   C  Southampton  yes  False          0
  4  man      True  NaN  Southampton  no  True          2
```

Applying One Hot Encoding technique

```
[17]: df_encoded = pd.get_dummies(df, columns=["sex", "embarked"], drop_first=True)
df_encoded.head()
```

```
[17]:    survived pclass   age  sibsp  parch      fare  class  who  adult_male \
  0        0      3  22.0     1      0    7.2500  Third  man      True
  1        1      1  38.0     1      0   71.2833  First woman     False
  2        1      3  26.0     0      0    7.9250  Third woman     False
```

```

3      1      1  35.0      1      0  53.1000  First woman      False
4      0      3  35.0      0      0  8.0500  Third man       True

   deck embark_town alive alone class_encoded sex_male embarked_Q \
0  NaN  Southampton  no  False          2     True    False
1   C   Cherbourg  yes  False          0    False    False
2  NaN  Southampton  yes   True          2    False    False
3   C  Southampton  yes  False          0    False    False
4  NaN  Southampton  no   True          2     True    False

embarked_S
0     True
1    False
2     True
3     True
4     True

```

3 Lab – 03: Exploratory Data Analysis (EDA)

Inspecting data using `head()`, `tail()`, `info()`, `describe()` methods

[18]: `df.head()`

```

[18]:    survived  pclass      sex   age  sibsp  parch      fare embarked  class \
0          0      3    male  22.0      1      0   7.2500      S  Third
1          1      1  female  38.0      1      0  71.2833      C  First
2          1      3  female  26.0      0      0   7.9250      S  Third
3          1      1  female  35.0      1      0  53.1000      S  First
4          0      3    male  35.0      0      0   8.0500      S  Third

      who  adult_male  deck  embark_town alive  alone  class_encoded
0  man      True  NaN  Southampton  no  False          2
1  woman    False   C   Cherbourg  yes  False          0
2  woman    False  NaN  Southampton  yes   True          2
3  woman    False   C  Southampton  yes  False          0
4  man      True  NaN  Southampton  no   True          2

```

[19]: `df.tail()`

```

[19]:    survived  pclass      sex   age  sibsp  parch      fare embarked  class \
886      0      2    male  27.0      0      0  13.00      S  Second
887      1      1  female  19.0      0      0  30.00      S  First
888      0      3  female  28.0      1      2  23.45      S  Third
889      1      1    male  26.0      0      0  30.00      C  First
890      0      3    male  32.0      0      0   7.75      Q  Third

      who  adult_male  deck  embark_town alive  alone  class_encoded

```

```

886    man      True   NaN  Southampton    no   True      1
887  woman     False    B  Southampton   yes   True      0
888  woman     False   NaN  Southampton    no  False      2
889    man      True    C  Cherbourg   yes   True      0
890    man      True   NaN  Queenstown   no   True      2

```

[20]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 16 columns):
 #   Column        Non-Null Count  Dtype  
--- 
 0   survived      891 non-null   int64  
 1   pclass         891 non-null   int64  
 2   sex            891 non-null   object  
 3   age            891 non-null   float64 
 4   sibsp          891 non-null   int64  
 5   parch          891 non-null   int64  
 6   fare            891 non-null   float64 
 7   embarked       889 non-null   object  
 8   class          891 non-null   category
 9   who             891 non-null   object  
 10  adult_male     891 non-null   bool    
 11  deck            203 non-null   category
 12  embark_town    889 non-null   object  
 13  alive           891 non-null   object  
 14  alone           891 non-null   bool    
 15  class_encoded  891 non-null   int64  
dtypes: bool(2), category(2), float64(2), int64(5), object(5)
memory usage: 87.6+ KB

```

[21]: df.describe()

```

[21]:    survived    pclass      age      sibsp      parch      fare \
count  891.000000  891.000000  891.000000  891.000000  891.000000  891.000000
mean   0.383838  2.308642  29.361582  0.523008  0.381594  32.204208
std    0.486592  0.836071  13.019697  1.102743  0.806057  49.693429
min    0.000000  1.000000  0.420000  0.000000  0.000000  0.000000
25%    0.000000  2.000000  22.000000  0.000000  0.000000  7.910400
50%    0.000000  3.000000  28.000000  0.000000  0.000000  14.454200
75%    1.000000  3.000000  35.000000  1.000000  0.000000  31.000000
max    1.000000  3.000000  80.000000  8.000000  6.000000  512.329200

      class_encoded
count    891.000000
mean     1.308642
std      0.836071

```

```
min          0.000000
25%         1.000000
50%         2.000000
75%         2.000000
max         2.000000
```

Calculate mean, median, mode, standard deviation, variance, and correlation for numerical features.

```
[22]: df.mean(numeric_only=True)
```

```
[22]: survived      0.383838
pclass        2.308642
age           29.361582
sibsp         0.523008
parch         0.381594
fare          32.204208
adult_male    0.602694
alone         0.602694
class_encoded 1.308642
dtype: float64
```

```
[23]: df.median(numeric_only=True)
```

```
[23]: survived      0.0000
pclass        3.0000
age           28.0000
sibsp         0.0000
parch         0.0000
fare          14.4542
adult_male    1.0000
alone         1.0000
class_encoded 2.0000
dtype: float64
```

```
[24]: df.mode(numeric_only=True)
```

```
[24]:    survived  pclass   age  sibsp  parch  fare  adult_male  alone \
0          0       3     28.0     0       0    8.05      True     True
                                         class_encoded
0                           2
```

```
[25]: df.std(numeric_only=True)
```

```
[25]: survived      0.486592
pclass        0.836071
age           13.019697
```

```
sibsp          1.102743  
parch         0.806057  
fare          49.693429  
adult_male    0.489615  
alone         0.489615  
class_encoded 0.836071  
dtype: float64
```

```
[26]: df.var(numeric_only=True)
```

```
[26]: survived      0.236772  
pclass        0.699015  
age           169.512498  
sibsp         1.216043  
parch         0.649728  
fare          2469.436846  
adult_male    0.239723  
alone         0.239723  
class_encoded 0.699015  
dtype: float64
```

```
[27]: df.corr(numeric_only=True)
```

```
[27]:      survived   pclass      age     sibsp     parch     fare \\\nsurvived   1.000000 -0.338481 -0.064910 -0.035322  0.081629  0.257307\npclass     -0.338481  1.000000 -0.339898  0.083081  0.018443 -0.549500\nage        -0.064910 -0.339898  1.000000 -0.233296 -0.172482  0.096688\nsibsp      -0.035322  0.083081 -0.233296  1.000000  0.414838  0.159651\nparch      0.081629  0.018443 -0.172482  0.414838  1.000000  0.216225\nfare        0.257307 -0.549500  0.096688  0.159651  0.216225  1.000000\nadult_male -0.557080  0.094035  0.247704 -0.253586 -0.349943 -0.182024\nalone      -0.203367  0.135207  0.171647 -0.584471 -0.583398 -0.271832\nclass_encoded -0.338481  1.000000 -0.339898  0.083081  0.018443 -0.549500\n\n      adult_male   alone  class_encoded\nsurvived   -0.557080 -0.203367   -0.338481\npclass     0.094035  0.135207   1.000000\nage        0.247704  0.171647   -0.339898\nsibsp     -0.253586 -0.584471   0.083081\nparch     -0.349943 -0.583398   0.018443\nfare       -0.182024 -0.271832   -0.549500\nadult_male 1.000000  0.404744   0.094035\nalone      0.404744  1.000000   0.135207\nclass_encoded 0.094035  0.135207   1.000000
```


Identifying missing values

```
[28]: df.isnull().sum()
```

```
[28]: survived      0  
pclass          0  
sex             0  
age             0  
sibsp          0  
parch          0  
fare            0  
embarked        2  
class           0  
who             0  
adult_male      0  
deck            688  
embark_town     2  
alive           0  
alone           0  
class_encoded   0  
dtype: int64
```

Detecting Outliers using Boxplot

```
[29]: plt.boxplot(df["fare"])  
plt.title("Boxplot of Fare")  
plt.show()
```



Detecting Outliers using IQR

```
[30]: Q1 = df["fare"].quantile(0.25)
Q3 = df["fare"].quantile(0.75)
IQR = Q3 - Q1

lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR

outliers = df[(df["fare"] < lower) | (df["fare"] > upper)]
print("Number of outliers:", outliers.shape[0])
```

Number of outliers: 116

Handling Outliers (Capping Method)

```
[31]: df["fare"] = np.where(df["fare"] > upper, upper, df["fare"])
df["fare"] = np.where(df["fare"] < lower, lower, df["fare"])
```

```
[32]: df.head()
```


```
[32]:    survived  pclass      sex   age  sibsp  parch     fare embarked  class \
0          0       3    male  22.0      1      0    7.2500      S  Third
1          1       1  female  38.0      1      0   65.6344      C  First
2          1       3  female  26.0      0      0    7.9250      S  Third
3          1       1  female  35.0      1      0   53.1000      S  First
4          0       3    male  35.0      0      0    8.0500      S  Third

      who  adult_male  deck  embark_town  alive  alone  class_encoded
0  man        True   NaN  Southampton   no  False           2
1 woman       False    C  Cherbourg   yes  False           0
2 woman       False   NaN  Southampton   yes  True            2
3 woman       False    C  Southampton   yes  False           0
4 man        True   NaN  Southampton   no  True            2
```

Plotting histograms, scatter plots, and bar charts using matplotlib or seaborn.


Histogram

```
[33]: plt.hist(df["age"], bins=20)
plt.title("Age Distribution")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.show()
```




Scatter Plot

```
[34]: plt.scatter(df["age"], df["fare"])
plt.xlabel("Age")
plt.ylabel("Fare")
plt.title("Age vs Fare")
plt.show()
```




Bar Plot

```
[35]: sns.countplot(x="survived", data=df)
plt.title("Survival Count")
plt.show()
```



Pairplot

```
[36]: sns.pairplot(df[["age", "fare", "pclass", "survived"]])  
plt.show()
```



Categorical Variable Analysis

```
[37]: df["sex"].value_counts()
```

```
[37]: sex
male      577
female    314
Name: count, dtype: int64
```


```
[38]: df["class"].value_counts()
```

```
[38]: class
Third     491
```

```
First      216
Second     184
Name: count, dtype: int64
```


Bar Plot for categorical variables

```
[39]: sns.countplot(x="sex", hue="survived", data=df)
plt.title("Survival by Gender")
plt.show()
```




Numerical Feature Distribution (Density Plot)

```
[40]: sns.kdeplot(df["age"], fill=True)
plt.title("Age Density Plot")
plt.show()
```



Boxplot for Numerical Feature

```
[41]: sns.boxplot(x="survived", y="fare", data=df)
plt.title("Fare Distribution by Survival")
plt.show()
```



4 Lab – 04: Linear Regression

Loading Advertising Dataset from local device

```
[42]: df = pd.read_csv("archive/Advertising.csv")
df.head()
```

```
[42]:   Unnamed: 0      TV  Radio  Newspaper  Sales
0            1    230.1    37.8      69.2   22.1
1            2     44.5    39.3      45.1   10.4
2            3     17.2    45.9      69.3    9.3
3            4    151.5    41.3      58.5   18.5
4            5    180.8    10.8      58.4   12.9
```

```
[43]: df.drop(columns='Unnamed: 0', inplace=True)
```

```
[44]: df.head()
```

```
[44]:      TV  Radio  Newspaper  Sales
0    230.1    37.8      69.2    22.1
1     44.5    39.3      45.1    10.4
2     17.2    45.9      69.3     9.3
3    151.5    41.3      58.5    18.5
4    180.8    10.8      58.4    12.9
```

For simple linear regression, taking single feature TV and target feature Sales

```
[45]: X = df[['TV']]
y = df['Sales']
```

Splitting dataset in train and test

```
[46]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

Training Model

```
[47]: from sklearn.linear_model import LinearRegression
model_simple = LinearRegression()
model_simple.fit(X_train, y_train)
```


```
[47]: LinearRegression()
```

Predicting

```
[48]: y_pred_simple = model_simple.predict(X_test)
```

Visualize Regression Line

```
[49]: plt.scatter(X_test, y_test)
plt.plot(X_test, y_pred_simple, color="red")
plt.xlabel("TV Budget")
plt.ylabel("Sales")
plt.title("Simple Linear Regression")
plt.show()
```



Evaluating Model

```
[50]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

print("MAE:", mean_absolute_error(y_test, y_pred_simple))
print("MSE:", mean_squared_error(y_test, y_pred_simple))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_simple)))
print("R2:", r2_score(y_test, y_pred_simple))
```

MAE: 2.444420003751042

MSE: 10.204654118800956

RMSE: 3.194472431998898

R2: 0.6766954295627077

For Multiple Linear regression

```
[51]: X_multi = df[["TV", "Radio", "Newspaper"]]
y_multi = df["Sales"]
```

Splitting data

```
[52]: X_train_m, X_test_m, y_train_m, y_test_m = train_test_split(X_multi, y_multi,  
    ↪test_size=0.2, random_state=42)
```

Train model

```
[53]: model_multi = LinearRegression()  
model_multi.fit(X_train_m, y_train_m)
```

```
[53]: LinearRegression()
```

Prediction

```
[54]: y_pred_multi = model_multi.predict(X_test_m)
```

Evaluate Multiple Model

```
[55]: print("MAE:", mean_absolute_error(y_test_m, y_pred_multi))  
print("MSE:", mean_squared_error(y_test_m, y_pred_multi))  
print("RMSE:", np.sqrt(mean_squared_error(y_test_m, y_pred_multi)))  
print("R2:", r2_score(y_test_m, y_pred_multi))
```

MAE: 1.4607567168117606


MSE: 3.1740973539761046

RMSE: 1.7815996615334502

R2: 0.899438024100912

Residual Plot

```
[56]: residuals = y_test_m - y_pred_multi  
plt.scatter(y_pred_multi, residuals)  
plt.axhline(y=0, color="red")  
plt.xlabel("Predicted Sales")  
plt.ylabel("Residuals")  
plt.title("Residual Plot")  
plt.show()
```



Predict on New Data

```
[57]: new_data = [[200, 40, 50]]
prediction = model_multi.predict(new_data)
print("Predicted Sales:", prediction)
```

Predicted Sales: [19.63082872]

5 Lab – 05: Logistic Regression

Load iris dataset

```
[58]: iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df[\"species\"] = iris.target
df[\"binary_species\"] = np.where(df[\"species\"]==0, 0, 1) # 0 = setosa, 1 = non-setosa
df.head()
```

```
[58]:    sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm) \
0           5.1          3.5            1.4            0.2
1           4.9          3.0            1.4            0.2
2           4.7          3.2            1.3            0.2
3           4.6          3.1            1.5            0.2
4           5.0          3.6            1.4            0.2

   species  binary_species
0         0             0
1         0             0
2         0             0
3         0             0
4         0             0
```

Splitting the data

```
[59]: X = df[iris.feature_names]
y = df["binary_species"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

Train Logistic model

```
[60]: from sklearn.linear_model import LogisticRegression

log_model = LogisticRegression()
log_model.fit(X_train, y_train)
```

[60]: LogisticRegression()

Making Prediction

```
[61]: y_pred = log_model.predict(X_test)
```

Accuracy Score

```
[62]: from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 1.0

Confusion Matrix & Classification Report

```
[63]: from sklearn.metrics import confusion_matrix, classification_report

cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", cm)
cr = classification_report(y_test, y_pred)
```

```
print("Classification Report:\n", cr)
```

Confusion Matrix:

```
[[10  0]
 [ 0 20]]
```


Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	20
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[64]: from sklearn.metrics import roc_curve, roc_auc_score
y_prob = log_model.predict_proba(X_test)[:,1] # probability for class 1

fpr, tpr, thresholds = roc_curve(y_test, y_prob)
auc_score = roc_auc_score(y_test, y_prob)

plt.plot(fpr, tpr, color="blue", label=f"AUC = {auc_score:.2f}")
plt.plot([0,1], [0,1], color="red", linestyle="--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```



Multi Class Classification

```
[65]: X_multi = df[iris.feature_names]
y_multi = df["species"]

X_train_m, X_test_m, y_train_m, y_test_m = train_test_split(
    X_multi, y_multi, test_size=0.2, random_state=42
)

log_model_multi = LogisticRegression(multi_class="ovr", max_iter=200)
log_model_multi.fit(X_train_m, y_train_m)

y_pred_m = log_model_multi.predict(X_test_m)
```

Multiclass Evaluation

```
[66]: print("Accuracy:", accuracy_score(y_test_m, y_pred_m))
print(classification_report(y_test_m, y_pred_m))
```

```
Accuracy: 0.9666666666666667
      precision    recall  f1-score   support

```

0	1.00	1.00	1.00	10
1	1.00	0.89	0.94	9
2	0.92	1.00	0.96	11
accuracy			0.97	30
macro avg	0.97	0.96	0.97	30
weighted avg	0.97	0.97	0.97	30

Predicting New Data

```
[67]: # Example new samples (sepal length, sepal width, petal length, petal width)
new_samples = [[5.1, 3.5, 1.4, 0.2],      # likely setosa
               [6.0, 2.9, 4.5, 1.5],    # likely versicolor
               [6.9, 3.1, 5.4, 2.1]]   # likely virginica

predicted_classes = log_model_multi.predict(new_samples)
predicted_names = [iris.target_names[i] for i in predicted_classes]

predicted_names
```

[67]: [np.str_('setosa'), np.str_('versicolor'), np.str_('virginica')]

6 Lab – 06: Decision Trees

```
[68]: df = df.drop(columns=['binary_species'])
df.head()
```

[68]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	species
0	0
1	0
2	0
3	0
4	0

```
[69]: from sklearn.tree import DecisionTreeClassifier, plot_tree
```

Splitting the data

```
[70]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42)
```


Train Decision Tree

```
[71]: dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)
```

```
[71]: DecisionTreeClassifier(random_state=42)
```

Visualize Tree

```
[72]: plt.figure(figsize=(12,8))
plot_tree(
    dt_model,
    feature_names=iris.feature_names,
    class_names=iris.target_names,
    filled=True
)
plt.show()
```



Predict

```
[73]: y_pred = dt_model.predict(X_test)
```

Performance Evaluation

```
[74]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred, average='weighted'))
print("Recall:", recall_score(y_test, y_pred, average='weighted'))
print("F1 Score:", f1_score(y_test, y_pred, average='weighted'))

print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))
```

Accuracy: 1.0
Precision: 1.0
Recall: 1.0
F1 Score: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	20
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Gini Vs Entropy

```
[75]: dt_gini = DecisionTreeClassifier(criterion="gini", random_state=42)
dt_gini.fit(X_train, y_train)

dt_entropy = DecisionTreeClassifier(criterion="entropy", random_state=42)
dt_entropy.fit(X_train, y_train)
```

```
[75]: DecisionTreeClassifier(criterion='entropy', random_state=42)
```

Compare Accuracy

```
[76]: print("Gini Accuracy:", accuracy_score(y_test, dt_gini.predict(X_test)))

print("Entropy Accuracy:", accuracy_score(y_test, dt_entropy.predict(X_test)))
# Usually same for iris dataset
```

Gini Accuracy: 1.0
Entropy Accuracy: 1.0

Pruned Tree

```
[77]: dt_pruned = DecisionTreeClassifier(  
        max_depth=3,  
        min_samples_split=4,  
        min_samples_leaf=2,  
        random_state=42  
)  
  
dt_pruned.fit(X_train, y_train)
```

```
[77]: DecisionTreeClassifier(max_depth=3, min_samples_leaf=2, min_samples_split=4,  
                             random_state=42)
```

Compare Accuracy

```
[78]: print("Unpruned Accuracy:", accuracy_score(y_test, dt_model.predict(X_test)))  
print("Pruned Accuracy:", accuracy_score(y_test, dt_pruned.predict(X_test)))
```

```
Unpruned Accuracy: 1.0  
Pruned Accuracy: 1.0
```

Tree Complexity

```
[79]: print("Unpruned depth:", dt_model.get_depth())  
print("Pruned depth:", dt_pruned.get_depth())
```

```
Unpruned depth: 1  
Pruned depth: 1
```

Predicting in new samples

```
[80]: new_samples = [  
        [5.1, 3.5, 1.4, 0.2],    # likely setosa  
        [6.0, 2.9, 4.5, 1.5],    # likely versicolor  
        [6.9, 3.1, 5.4, 2.1]    # likely virginica  
    ]  
  
predictions = dt_pruned.predict(new_samples)  
  
predicted_names = [iris.target_names[i] for i in predictions]  
  
predicted_names
```

```
[80]: [np.str_('setosa'), np.str_('versicolor'), np.str_('versicolor')]
```

7 Lab – 07: Random Forests and Ensemble Methods

Load Titanic dataset

```
[81]: df = sns.load_dataset("titanic")
df = df[["survived", "pclass", "sex", "age", "sibsp", "parch", "fare", "embarked"]]
df.head()
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S

Basic Preprocessin

```
[82]: # Fill missing values
df["age"].fillna(df["age"].median(), inplace=True)
df["embarked"].fillna(df["embarked"].mode()[0], inplace=True)

# One-hot encoding
df = pd.get_dummies(df, columns=["sex", "embarked"], drop_first=True)
```

Train-Test Split

```
[83]: X = df.drop("survived", axis=1)
y = df["survived"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

Train Random Forest Classifier

```
[84]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

rf_model = RandomForestClassifier(n_estimators=100, random_state=42)

rf_model.fit(X_train, y_train)

y_pred = rf_model.predict(X_test)

print("Random Forest Accuracy:", accuracy_score(y_test, y_pred))
```

Random Forest Accuracy: 0.8212290502793296

Important feature analysis

```
[85]: importances = rf_model.feature_importances_

importance_df = pd.DataFrame({
    "Feature": X.columns,
    "Importance": importances}
```


```
}).sort_values(by="Importance", ascending=False)

importance_df.head()
```

```
[85]:    Feature  Importance
5   sex_male    0.273316
4      fare     0.272058
1      age      0.252745
0   pclass     0.078616
2    sibsp     0.052192
```

Visualizing Important feature

```
[86]: plt.figure(figsize=(8,5))
plt.barh(importance_df["Feature"], importance_df["Importance"])
plt.gca().invert_yaxis()
plt.title("Feature Importance (Titanic)")
plt.show()
```



Hyperparameter Tuning

Model_1: small

```
[87]: rf_small = RandomForestClassifier(n_estimators=50, max_depth=3,random_state=42)

rf_small.fit(X_train, y_train)
```

```
print("Small Forest Accuracy:", accuracy_score(y_test, rf_small.  
predict(X_test)))
```

Small Forest Accuracy: 0.8044692737430168

Model_2: Large

```
[88]: rf_large = RandomForestClassifier(  
        n_estimators=200,  
        max_depth=None,  
        min_samples_split=2,  
        min_samples_leaf=1,  
        random_state=42  
)  
  
rf_large.fit(X_train, y_train)  
  
print("Large Forest Accuracy:", accuracy_score(y_test, rf_large.  
predict(X_test)))
```

Large Forest Accuracy: 0.8100558659217877

Observation - More trees more stable predictions - Lower depth less overfitting - Very deep trees may overfit

7.1 Random Forest Regressor

```
[89]: from sklearn.datasets import fetch_california_housing  
from sklearn.ensemble import RandomForestRegressor
```

```
[90]: data = fetch_california_housing()  
  
X = data.data  
y = data.target  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)  
  
rf_reg = RandomForestRegressor(n_estimators=100, random_state=42)  
rf_reg.fit(X_train, y_train)  
  
y_pred = rf_reg.predict(X_test)  
  
print("MSE:", mean_squared_error(y_test, y_pred))  
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))  
print("R2:", r2_score(y_test, y_pred))
```

MSE: 0.2553684927247781

RMSE: 0.5053399773665033

R2: 0.8051230593157366

8 Lab – 08: Support Vector Machine (SVM)

Loading Iris Data

```
[91]: iris = load_iris()
X = iris.data
y = iris.target

# Keep only class 0 and 1
X = X[y != 2]
y = y[y != 2]

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42
)
```

Train Linear SVM

```
[92]: from sklearn.svm import SVC
svm_linear = SVC(kernel="linear")
svm_linear.fit(X_train, y_train)

y_pred = svm_linear.predict(X_test)

print("Linear SVM Accuracy:", accuracy_score(y_test, y_pred))
```

Linear SVM Accuracy: 1.0

Using Different Kernels

```
[93]: kernels = ["linear", "poly", "rbf"]

for k in kernels:
    model = SVC(kernel=k)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"{k} Kernel Accuracy:", accuracy_score(y_test, y_pred))
```

linear Kernel Accuracy: 1.0
poly Kernel Accuracy: 1.0
rbf Kernel Accuracy: 1.0

Hyperparameter Tuning

```
[94]: for c in [0.1, 1, 10, 100]:
    model = SVC(kernel="linear", C=c)
```

```

model.fit(X_train, y_train)
print(f"C={c} Accuracy:",
      accuracy_score(y_test, model.predict(X_test)))

```

C=0.1 Accuracy: 1.0
C=1 Accuracy: 1.0
C=10 Accuracy: 1.0
C=100 Accuracy: 1.0

Effect of Gamma (RBF)

```
[95]: for g in [0.01, 0.1, 1, 10]:
    model = SVC(kernel="rbf", gamma=g)
    model.fit(X_train, y_train)
    print(f"gamma={g} Accuracy:",
          accuracy_score(y_test, model.predict(X_test)))
```

gamma=0.01 Accuracy: 1.0
gamma=0.1 Accuracy: 1.0
gamma=1 Accuracy: 1.0
gamma=10 Accuracy: 0.95

Multi-class Classification

```
[96]: # Full dataset
X_full = iris.data
y_full = iris.target

X_train, X_test, y_train, y_test = train_test_split(
    X_full, y_full,
    test_size=0.2,
    random_state=42
)

svm_multi = SVC(kernel="rbf")
svm_multi.fit(X_train, y_train)

print("Multiclass Accuracy:",
      accuracy_score(y_test, svm_multi.predict(X_test)))
```

Multiclass Accuracy: 1.0

Visualizing Decision Boundary (2D)

```
[97]: X_2d = X[:, :2]    # first two features
y_2d = y

svm_vis = SVC(kernel="linear")
svm_vis.fit(X_2d, y_2d)
```

```
[97]: SVC(kernel='linear')
```

```
[98]: def plot_svm(model, X, y):
    plt.scatter(X[:,0], X[:,1], c=y, cmap="coolwarm")

    ax = plt.gca()
    xlim = ax.get_xlim()
    ylim = ax.get_ylim()


    xx = np.linspace(xlim[0], xlim[1], 30)
    yy = np.linspace(ylim[0], ylim[1], 30)
    YY, XX = np.meshgrid(yy, xx)
    xy = np.vstack([XX.ravel(), YY.ravel()]).T
    Z = model.decision_function(xy).reshape(XX.shape)

    ax.contour(XX, YY, Z, levels=[0], colors="black")

    # Support vectors
    ax.scatter(model.support_vectors_[:,0],
               model.support_vectors_[:,1],
               s=100, facecolors="none",
               edgecolors="k")

    plt.title("SVM Decision Boundary")
    plt.show()

plot_svm(svm_vis, X_2d, y_2d)
```



9 Lab – 09: K-Nearest Neighbors (KNN)

```
[99]: iris = load_iris()

X = iris.data
y = iris.target
```

Feature Scaling

```
[100]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Train test split

```
[101]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)
```

KNN classification

```
[102]: from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

Accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Compare Different K Values

```
[103]: for k in [1, 3, 5, 7, 9]:
    model = KNeighborsClassifier(n_neighbors=k)
    model.fit(X_train, y_train)
    print(f"K={k} Accuracy:", accuracy_score(y_test, model.predict(X_test)))
```

K=1 Accuracy: 0.9666666666666667
K=3 Accuracy: 1.0
K=5 Accuracy: 1.0
K=7 Accuracy: 1.0
K=9 Accuracy: 1.0

Compare Distance Metrics

```
[104]: for metric in ["euclidean", "manhattan", "minkowski"]:
    model = KNeighborsClassifier(n_neighbors=5, metric=metric)
    model.fit(X_train, y_train)
    print(metric, "Accuracy:", accuracy_score(y_test, model.predict(X_test)))
```

euclidean Accuracy: 1.0
manhattan Accuracy: 1.0
minkowski Accuracy: 1.0

Decision Boundary (2D Visualization)

```
[105]: X_2d = X_scaled[:, :2] # first 2 features
```

```
knn_2d = KNeighborsClassifier(n_neighbors=5)
knn_2d.fit(X_2d, y)
```


[105]: KNeighborsClassifier()

```
[106]: h = 0.02
x_min, x_max = X_2d[:, 0].min() - 1, X_2d[:, 0].max() + 1
y_min, y_max = X_2d[:, 1].min() - 1, X_2d[:, 1].max() + 1

xx, yy = np.meshgrid(
    np.arange(x_min, x_max, h),
    np.arange(y_min, y_max, h)
)

Z = knn_2d.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.contourf(xx, yy, Z, alpha=0.3)
plt.scatter(X_2d[:, 0], X_2d[:, 1], c=y)
plt.title("KNN Decision Boundary (Iris)")
plt.show()
```



10 Lab – 10: Unsupervised Learning – Clustering

```
[107]: from sklearn.cluster import KMeans, DBSCAN
from sklearn.metrics import silhouette_score, davies_bouldin_score
from scipy.cluster.hierarchy import dendrogram, linkage
```

Load Dataset (Without Labels)

```
[108]: iris = load_iris()
X = iris.data    # only features


scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

K-Means Clustering

```
[109]: kmeans = KMeans(n_clusters=3, random_state=42)
clusters = kmeans.fit_predict(X_scaled)
```

Visualize Clusters

```
[110]: plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=clusters)
plt.title("K-Means Clustering")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()
```




Elbow Method (Optimal K)

```
[111]: wcss = []

for k in range(1, 11):
    model = KMeans(n_clusters=k, random_state=42)
    model.fit(X_scaled)
    wcss.append(model.inertia_)

plt.plot(range(1, 11), wcss)
plt.xlabel("Number of Clusters")
plt.ylabel("WCSS")
plt.title("Elbow Method")
plt.show()
```




Hierarchical Clustering

```
linked = linkage(X_scaled, method='ward')
plt.figure(figsize=(8,5)) dendrogram(linked) plt.title("Hierarchical Clustering Dendrogram")
plt.show()
```

DBSCAN Clustering

```
[112]: dbSCAN = DBSCAN(eps=0.8, min_samples=5)
db_clusters = dbSCAN.fit_predict(X_scaled)

plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=db_clusters)
plt.title("DBSCAN Clustering")
plt.show()
```



Silhouette Score

```
[113]: print("KMeans Silhouette:", silhouette_score(X_scaled, clusters))
print("DBSCAN Silhouette:", silhouette_score(X_scaled, db_clusters))
```

KMeans Silhouette: 0.4798814508199817
DBSCAN Silhouette: 0.5216965052515835

Davies-Bouldin Index

```
[114]: print("KMeans DB Index:", davies_bouldin_score(X_scaled, clusters))
print("DBSCAN DB Index:", davies_bouldin_score(X_scaled, db_clusters))
```

KMeans DB Index: 0.7893630242997912
DBSCAN DB Index: 1.9432005358011466

11 Lab – 11: Dimensionality Reduction

```
[115]: from sklearn.datasets import load_digits
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
```

Load High-Dimensional Dataset

```
[116]: digits = load_digits()  
X = digits.data      # 64 features  
y = digits.target    # 0-9 classes  
  
print("Original shape:", X.shape)
```

Original shape: (1797, 64)

Standardize Data

```
[117]: scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```


Apply PCA (2 Components for Visualization)

```
[118]: pca = PCA(n_components=2)  
X_pca = pca.fit_transform(X_scaled)  
  
print("Reduced shape:", X_pca.shape)
```

Reduced shape: (1797, 2)

Visualize First Two Principal Components

```
[119]: plt.figure(figsize=(6,5))  
plt.scatter(X_pca[:,0], X_pca[:,1], c=y)  
plt.title("PCA - First Two Principal Components")  
plt.xlabel("PC1")  
plt.ylabel("PC2")  
plt.show()
```




Variance

```
[120]: pca_full = PCA()
pca_full.fit(X_scaled)

cumulative_variance = np.cumsum(pca_full.explained_variance_ratio_)

plt.plot(cumulative_variance)
plt.xlabel("Number of Components")
plt.ylabel("Cumulative Explained Variance")
plt.title("Variance Explained by PCA")
plt.show()
```



Number of Components for 95% Variance

```
[121]: n_components_95 = np.argmax(cumulative_variance >= 0.95) + 1
print("Components for 95% variance:", n_components_95)
```

Components for 95% variance: 40

Reduce to 95% Variance

```
[122]: pca_95 = PCA(n_components=n_components_95)
X_reduced = pca_95.fit_transform(X_scaled)
```

Train Classifier (Logistic Regression)

```
[123]: X_train, X_test, y_train, y_test = train_test_split(
    X_reduced, y,
    test_size=0.2,
    random_state=42
)

clf = LogisticRegression(max_iter=2000)
clf.fit(X_train, y_train)
```

```

y_pred = clf.predict(X_test)

print("Accuracy with PCA:", accuracy_score(y_test, y_pred))

```

Accuracy with PCA: 0.9611111111111111


t-SNE Visualization

```

[124]: tsne = TSNE(n_components=2, random_state=42)
X_tsne = tsne.fit_transform(X_scaled)

plt.figure(figsize=(6,5))
plt.scatter(X_tsne[:,0], X_tsne[:,1], c=y)
plt.title("t-SNE Visualization")
plt.show()

```



Observations: - t-SNE gives better visual cluster separation. - PCA is suitable for feature reduction before ML. - t-SNE is mainly for visualization.

12 Lab – 12: Model Evaluation & Cross-Validation

```
[125]: from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import KFold, StratifiedKFold, cross_val_score, GridSearchCV
```

Load and Scale the dataset

```
[126]: data = load_breast_cancer()
X = data.data
y = data.target

scaler = StandardScaler()
X = scaler.fit_transform(X)
```

Train-Test Split Evaluation

```
[127]: X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42
)

model = LogisticRegression(max_iter=2000)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

Accuracy: 0.9736842105263158

Precision: 0.9722222222222222

Recall: 0.9859154929577465

F1 Score: 0.9790209790209791

K-Fold Cross-Validation

```
[128]: kf = KFold(n_splits=5, shuffle=True, random_state=42)

cv_scores = cross_val_score(model, X, y, cv=kf)

print("KFold Accuracy Scores:", cv_scores)
print("Average Accuracy:", np.mean(cv_scores))
```

KFold Accuracy Scores: [0.97368421 0.98245614 0.96491228 0.99122807 0.97345133]

Average Accuracy: 0.9771464058376029

Stratified K-Fold Cross-Validation

```
[129]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

strat_scores = cross_val_score(model, X, y, cv=skf)

print("StratifiedKFold Scores:", strat_scores)
print("Average Accuracy:", np.mean(strat_scores))
```

StratifiedKFold Scores: [0.97368421 0.94736842 0.96491228 0.99122807 0.99115044]
 Average Accuracy: 0.9736686849868033

Hyperparameter Tuning with Grid Search

```
[130]: param_grid = {
    'C': [0.1, 1, 10],
    'gamma': ['scale', 0.01, 0.001],
    'kernel': ['rbf']
}

grid = GridSearchCV(
    SVC(probability=True),
    param_grid,
    cv=5,
    scoring='accuracy'
)

grid.fit(X_train, y_train)

print("Best Parameters:", grid.best_params_)
print("Best CV Score:", grid.best_score_)
```

Best Parameters: {'C': 1, 'gamma': 'scale', 'kernel': 'rbf'}
 Best CV Score: 0.9736263736263737

Evaluate on test set:

```
[131]: best_model = grid.best_estimator_
y_pred = best_model.predict(X_test)

print("Test Accuracy:", accuracy_score(y_test, y_pred))
```

Test Accuracy: 0.9736842105263158

ROC Curve and AUC


```
[132]: y_prob = best_model.predict_proba(X_test)[:, 1]

fpr, tpr, thresholds = roc_curve(y_test, y_prob)
auc_score = roc_auc_score(y_test, y_prob)

plt.plot(fpr, tpr)
```

```
plt.plot([0,1], [0,1], linestyle='--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.show()
```

```
print("AUC Score:", auc_score)
```



AUC Score: 0.99737962659679

[]: